

# Exploiting Generational Garbage Collection

## Using Data Remnants to Improve Memory Analysis and Digital Forensics

**Adam Pridgen**<sup>1</sup>

<sup>1</sup>Rice University, Houston, TX, USA

January 18, 2017



1

## Introduction

- Motivation
- Contributions
- Background

2

## Supporting Work

- STAAF: Scaling Android Application Analysis
- Radare Java Static Analysis

3

## Present but unreachable

- Introduction
- Background
- Problem
- Approach
- Results
- Conclusions

4

## Picking up the Trash

- Introduction
- Problem
- Approach
- Evaluation
- Conclusions

5

## Conclusions

6

## References

# Malware Compromises Happen

RICE



**How did it happen?**

**Who did it?**

**Was it targeted?**

**Why did this happen?**

**What was lost?**

# Motivation (1)




Figure: Stuxnet riding over an airgap into an ICS network [1].

# Motivation (2)

- Threats actors looking to penetrate hard-targets
  - Must research and innovate on existing methods
  - Exploit technology blind spots and implicit trust
  - Looking for *knowledge-gaps*
  - Targetting technologies that are widely deployed
  - Exploring exploitation in all dimensions

# Motivation (3)

- **Managed Runtimes**

- Widely used but not well understood
- Runtimes are complex and evolve over time
- Backwards compatibility retained
- Widely deployed runs on multiple platforms
- Updates may not be feasible

# Contributions

- Developed tools for Java class and archive analysis
- Established the feasibility of recovering artifacts
- Created an approach for recovering managed objects
- Developed a prototype targeting the HotSpot JVM



# Attacker Techniques



Figure: Overview of attacker tactics.

# Documented Java Malware and Attacks

- Malware and Backdoors
  - **Criminal**: Adwind, JBot, etc. [2, 3, 4]
  - **Espionage**: PackRat and JavaFog [5, 6]
- Threat actors employing Java
  - **Phishing** [7, 4]
  - **Waterholing** [7, 4]
- Common Vulnerabilities and Exposure
  - **Hotspot JVM**:  $\approx 34$  since 2010 [8]
  - **Java and frameworks**:  $\approx 1510$  since 1999

# Digital Forensics Overview

## Acquisition

### Capture Artifacts

1. System Logs
2. Memory Dumps
3. Network Traffic
4. Network Access Logs
5. Application Logs
6. Disk drive images
7. ...

## Analysis

### Identify Relevant Details

1. Assemble log events
2. Disk drive modification
3. Memory analysis
4. Network activity
5. Determine movements
6. Enumerate access
7. ...

## Report

### Enable actionable recovery and mitigation steps

1. How?
2. Why?
3. When?
4. Duration?
5. Mitigation?
6. Litigation or Prosecution?
7. ...

# Related Work: Managed Runtime Analysis

- Viega explains the insecurity of managed runtimes [9]
- Chow et al. solve secure deallocation on Unix [10, 11]
- CleanOS: Objects encrypted using a shared key [12]
- Anikeev et al. focuses on Android's collector [13]
- Li shows RSA keys are retrievable in Python [14]

# Related Work: Memory Analysis

- Rekall and Volatility analysis frameworks [15, 16]
- Identifying datastructures
  - Lin et al. perform automatic RE [17]
  - Lin et al. use graph-based signatures [18]
  - Dolan et al. focus on kernel structures [19]
- Android memory forensics [20, 21, 22, 23, 24, 25]
- Data carving
  - Richard developed Scalpel File Carver [26]
  - Beverly et al. focus on network packets [27]
  - Hand et al. extract binaries from memory [28]

# Contents

1

Introduction

● Motivation

● Contributions

● Background

2

**Supporting Work**

● STAAF: Scaling Android Application Analysis

● Radare Java Static Analysis

3

**Present but unreachable**

● Introduction

● Background

● Problem

● Approach

● Results

● Conclusions

4

**Picking up the Trash**

● Introduction

● Problem

● Approach

● Evaluation

● Conclusions

5

Conclusions

6

References

# STAAF

## STAAF: Scaling Android Application Analysis with a Modular Framework

**Ryan W. Smith**<sup>1</sup>   **Adam Pridgen**<sup>1</sup>

<sup>1</sup>Praetorian, Austin, TX, USA

<sup>2</sup>Rice University, Houston, TX, USA

Hawaii International Conference on System Sciences, 2012



# Problem

- Engineering scalable program analysis
- Off-market Android stores were contained malware
- Android analysis tools fail to scale alone
- Developed an approach pipelining analysis



- **Similar approach used to measure latent secrets**
- Emphasizes scaling analysis horizontally
- Creates a pipeline for pre- and post- analysis
- Efficiently localizes analysis results in a database

# Reversing Java (Malware) with Radare

## STAAF: Scaling Android Application Analysis with a Modular Framework

**Adam Pridgen**<sup>1</sup>

<sup>1</sup>Rice University, Houston, TX, USA

InfoSec Southwest, 2014



# Problem

- Malware obfuscation would throw-off analysis
- Tools were built on Java
- Tools overlooked corner cases
- None of the tools allowed low-level manipulation

# Java Malware Analysis Overview

- **Eclipse IDE**: development environment for debugging
- **IDA Pro**: marked-up analysis with no low-level access
- **JD GUI**: decompiles code but cannot be corrected
- **Jython**: run Java in Python environment

# Results: Radare Extensions

- Low-level JAR and class file manipulation
- Analysis of class file artifacts
- Inject byte code for runtime analysis
- Rewrite symbolic links for hooking

1

Introduction

- Motivation
- Contributions
- Background

2

Supporting Work

- STAAF: Scaling Android Application Analysis
- Radare Java Static Analysis

3

**Present but unreachable**

- Introduction
- Background
- Problem
- Approach
- Results
- Conclusions

4

Picking up the Trash

- Introduction
- Problem
- Approach
- Evaluation
- Conclusions

5

Conclusions

6

References

# Present but unreachable

Reducing persistent latent secrets in HotSpot JVM  
***Best Paper**, Software Technology Track*

**Adam Pridgen**<sup>1</sup>   Simson L. Garfinkel<sup>2</sup>   Dan S. Wallach<sup>1</sup>

<sup>1</sup>Rice University, Houston, TX, USA

<sup>2</sup>George Mason University, Fairfax, VA, USA

Hawaii International Conference on System Sciences, 2017



# Introduction

- Java runtime uses automatic memory management
- Developers no longer control data lifetimes
- Sensitive data cannot be explicitly destroyed
- Multiple copies can be created



# Research Questions

- How many secrets are retained?
- Should we be concerned?
- Can we fix the problem (without vendor intervention)?
- Is our solution useful?

# Generational GC Heap Overview

- Tracing GC: Looking for *live* objects from a set of roots
- Heap engineered for expected object life-time
- GC promotes objects from one heap to the next one
  - **Eden Space** (short lived) → **Survivor Space**
  - **Survivor Space** → **Tenure Space** (long lived)




Figure: Typical generational heap layout.

# Other Factors Affecting Measurement

- GC algorithms and various collection conditions
- Internal JVM memory management system
- Interactions between JVM internals and program data
- Java Native Interface (not evaluated)

# Unmanaged Data Lifetime Overview

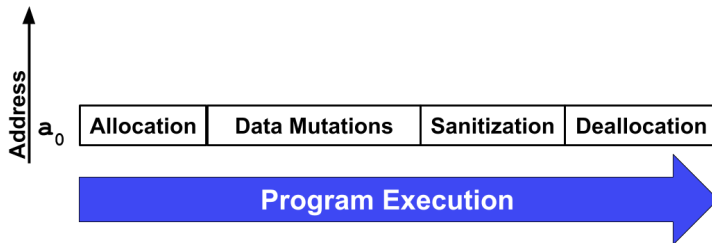


Figure: Example data lifetime in unmanaged memory.

# Managed Data Lifetime Overview

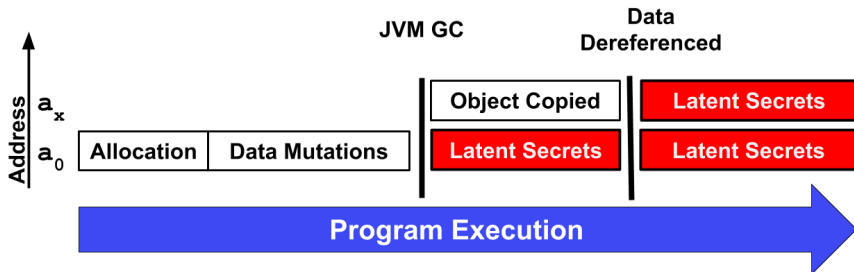


Figure: Example data lifetime in managed memory.

# Why is data being retained?

RICE

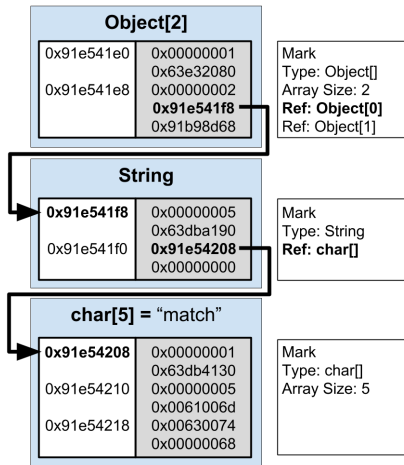


Figure: String[2] on the heap.

# Why is data being retained? (2)




Figure: String[0] is reassigned but the old value remains.

# Measuring Latent Secrets: Methodology

- Quantify data retention using TLS Keys
  - Vary memory pressure
  - Use well-known software examples
  - Vary heap size 512MiB-16GiB
- Modify HotSpot JVM to perform sanitization
- Re-evaluate data retention
- Measure the performance impacts



# Measuring Latent Secrets: TLS Clients

## Basic TLS Client

1. Wrap TLS socket
2. Manual HTTP communication
3. Rely on the Java Cryptography library

## Apache HTTP TLS Client

1. Library creates socket
2. Apache handles the communication
3. Rely on the Java Cryptography library

## Apache HTTP TLS Client with BouncyCastle

1. Library creates socket
2. Apache handles the communication
3. Rely on the BouncyCastle Cryptography library

# Measuring Latent Secrets: Memory Pressure

## High Memory Pressure

1. High Memory Contention
2. Consume up to 80%
3. 192 requests per running session (thread)

## Low Memory Pressure

1. Low Memory Contention
2. Consume up to 20%
3. 48 requests per running session (thread)

# Measuring Latent Secrets: Test Bench

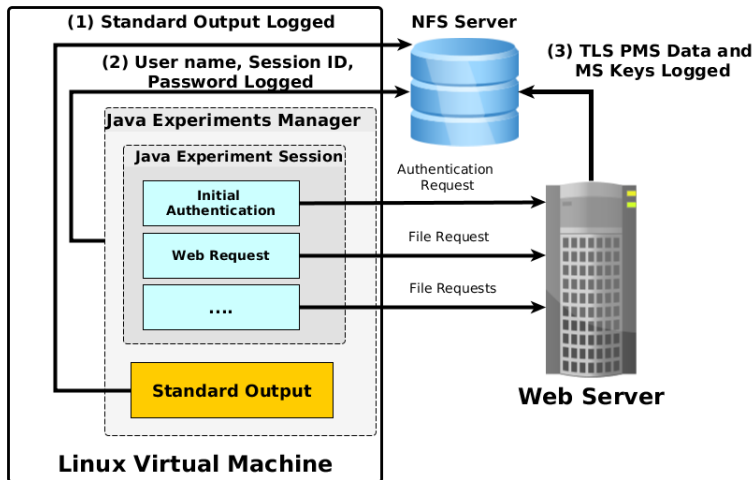


Figure: Overview of experiment and captured data.

# Measuring Latent Secrets: Data Processing

- Dump virtual machine system memory (e.g. RAM)
- Grep *RAM* for captured TLS key material
- Reconstruct the JVM process memory
- Grep *process memory* for TLS key material
- Reorder TLS sessions and count keys

# Reducing Latent Secrets

## Failed Approach

- Modify the Java Cryptography TLS Routines
  - Sanitize *out-of-scope* references
  - Explicit clean-up when sockets close or shutdown
- **Increased** the number of latent secrets

# Reducing Latent Secrets

## Successful Implementation

- Modify the JVM and GC algorithms
- Zero unused space after each collection
- Zero internally managed memory when deallocated

## Limitations

- Dangling references still prevent object's collection
- GC must occur on each heap space
- Sanitization may not be timely

# Reducing Latent Secrets


## Successful Implementation

- Modify the JVM and GC algorithms
- Zero unused space after each collection
- Zero internally managed memory when deallocated


## Limitations

- Dangling references still prevent object's collection
- GC must occur on each heap space
- Sanitization may not be timely

# Results - SerialGC HMP



(a) Socket Results




(b) Apache Results


Figure: TLS keys recovered from HMP clients.



# Results - SerialGC LMP




(a) Socket Results




(b) Apache Results

Figure: TLS keys recovered from LMP clients.

# Results - G1GC Sockets Client




(a) HMP Results




(b) LMP Results

Figure: TLS keys recovered from Socket clients using G1GC.

# Benchmarking Results



(a) tradebeans-Day Trader



(b) lusearch-Text Searching

Figure: Benchmarks show modifications reduced performance.

# Conclusions

- Quantified data retention in the HotSpot JVM
- Measured these secrets in a general manner
- Developed several strategies to reduce latent secrets
- **Data security** at the expense of **performance**

# Contents

1

Introduction

● Motivation

● Contributions

● Background

2

Supporting Work

● STAAF: Scaling Android Application Analysis

● Radare Java Static Analysis

3

Present but unreachable

● Introduction

● Background

● Problem

● Approach

● Results

● Conclusions

4

Picking up the Trash

● Introduction

● Problem

● Approach

● Evaluation

● Conclusions

5

Conclusions

6

References

# Picking up the trash

## Exploiting generational GC for memory analysis

**Adam Pridgen**<sup>1</sup>   **Simson L. Garfinkel**<sup>2</sup>   **Dan S. Wallach**<sup>1</sup>

<sup>1</sup>Rice University, Houston, TX, USA

<sup>2</sup>George Mason University, Fairfax, VA, USA

Digital Forensics Research Workshop, 2017




# Motivation

- Authors develop malware for managed runtimes
- Threat actors exploit vulnerable internet applications
- Managed runtimes retain artifacts [29]
- Digital forensics exploit this for evidence recovery

# Minding the Gap: Semantic Gaps


RICE






# Minding the Gap: OS View

RICE





# Minding the Gap: JVM Process View

RICE



# Minding the Gap: Java Thread View

RICE



# Attackers Advantages


- Tools and endpoint detection lack introspection
- Vulnerabilities exist in the *entire* software stack
- VMs are porous once DMA is achieved
- Complex machines limit *direct* analysis

# Analyst Advantages

- Separation of *code* and *data* limits attacker tricks
- Many artifacts for event reconstruction
- *Timelining* (e.g. artifacts ordered by creation)


# Recovering OOP Framework

RICE



# Recovering OOP Framework (2)

- Focuses recovery from x86 architecture
- Uses a minimal set of structure overlays
- Compatible with Linux and Windows OS



# Memory Capture and Reconstruction

- **Capture:** system memory (e.g. RAM) is dumped
- **Reconstruction:** Target process is found and page frames are reordered



# Extract Loaded Types

- Identify structures revealing loaded types
  - `SystemDictionary`: loaded classes
  - `SymbolTable`: loaded symbols)
  - `StringTable`: constants or long-lived strings
- Mine structures for the loaded data structures

# Extract Loaded Types (2)

- Look for invariant values
- Walk the hash tables
- Use constraints to control recovery
- Event ordering and timelining using addresses





# Locate Managed Memory with Logs

**Table:** The regular expression “`space.*used`” used in conjunction with `ffastrings` to determine the eden, survivor, and tenure generation spaces. Note `[...]` signifies omitted message content.

| GC Log Message     |                            |   |
|--------------------|----------------------------|---|
| Generational Space | Start and End of the Space |   |
| <b>eden space</b>  | <code>[...] used</code>    | <code>[0xa4800000, [...] 0xa4c50000)</code> |
| <b>from space</b>  | <code>[...] used</code>    | <code>[0xa4c50000, [...] 0xa4cd0000)</code> |
| <b>to space</b>    | <code>[...] used</code>    | <code>[0xa4cd0000, [...] 0xa4d50000)</code> |
| <b>the space</b>   | <code>[...] used</code>    | <code>[0xa9d50000, [...] 0xaa800000)</code> |

# Locating Managed Memory with Pointers

**Table:** Java object distribution in managed process memory (e.g. eden, survivor, and tenure spaces).

| Address Range         | Type Pointers | Unique Pointers | Pointer Occurrences Per Page (Y-axis: 0-64)  |
|-----------------------|---------------|-----------------|--|
| 0xa47ff000-0xa4c0f000 | 13261         | 266             |  |
| 0xa4c50000-0xa4c92000 | 129           | 28              |  |
| 0xa4cd0000-0xa4d50000 | 1121          | 79              |  |
| 0xa9d50000-0xaa000000 | 28810         | 661             |  |

## Locating Managed Memory with Pointers (2)

**Table:** Java address distribution in unmanaged process memory.

| Address Range         | Type<br>Pointers | Unique<br>Pointers | Pointer Occurrences<br>Per Page (Y-axis: 0-64) |
|-----------------------|------------------|--------------------|--|
| 0xa32de000-0xa3355000 | 1353             | 265                |  |
| 0xa33ce000-0xa349d000 | 2735             | 331                |  |
| 0xa349e000-0xa34f5000 | 609              | 122                |  |
| 0xa3600000-0xa3692000 | 362              | 360                |  |
| 0xc0001000-0xf7bfe000 | 11085            | 1211               |  |

# Enumerate and Extract Objects

- Scan managed heap for known-types
- Parse the object based on the report type
- Lift values for the object's fields

# Objects of Immediate Interest

**Java Threads**

**Sockets**

**Process Builders**

**Native Buffers**

**Streams**

**Child Processes**

**JAR Files**

**Buffers**

**JAR Entries**

# Malware Overview

- Adwind Malware Functionality
  - Obfuscation loads platform independent libraries
  - Modular design with plugins
  - Password protected with hard-coded IP address
  - Tested sample: beaconing, screen capture, etc.



# Malware Overview




Figure: Overview of the malware functionality for experiment.

# Reconstructing Socket Connections

| Object Address | Remote Connection | In/Out |   | Data (Up to 30 Bytes)          |
|----------------|-------------------|--------|---|--------------------------------|
| 0x91c779b8     | 10.18.120.18      | 48002  | ⇒ | Do something evil-48002!       |
| 0x91c7ead0     | 10.18.120.18      | 48003  | ⇒ | Do something evil-48003!       |
| 0x91c85b70     | 10.18.120.18      | 48002  | ⇐ | s3cr3t_d4t3_48002-00000000s3cr |
| 0x91c938d8     | 172.16.124.15     | 58860  | ⇒ | czNjcjN0X2Q0dDNfNDgwMDItMDAw   |
| 0x91c980d0     | 10.18.120.18      | 48003  | ⇐ | s3cr3t_d4t3_48003-00000000s3cr |
| 0x91ca5cb8     | 172.16.124.15     | 58860  | ⇒ | czNjcjN0X2Q0dDNfNDgwMDMtMDAw   |
| 0x91cbfef0     | 10.18.120.18      | 48004  | ⇒ | Do something evil-48004!       |
| 0x91cc7008     | 10.18.120.18      | 48005  | ⇒ | Do something evil-48005!       |
| 0x91ccdee8     | 10.18.120.18      | 48004  | ⇐ | s3cr3t_d4t3_48004-00000000s3cr |
| 0x91cdbad0     | 172.16.124.15     | 58860  | ⇒ | czNjcjN0X2Q0dDNfNDgwMDQtMDAw   |
| 0x91ce02c8     | 10.18.120.18      | 48005  | ⇐ | s3cr3t_d4t3_48005-00000000s3cr |
| 0x91cedeb0     | 172.16.124.15     | 58860  | ⇒ | czNjcjN0X2Q0dDNfNDgwMDUtMDAw   |

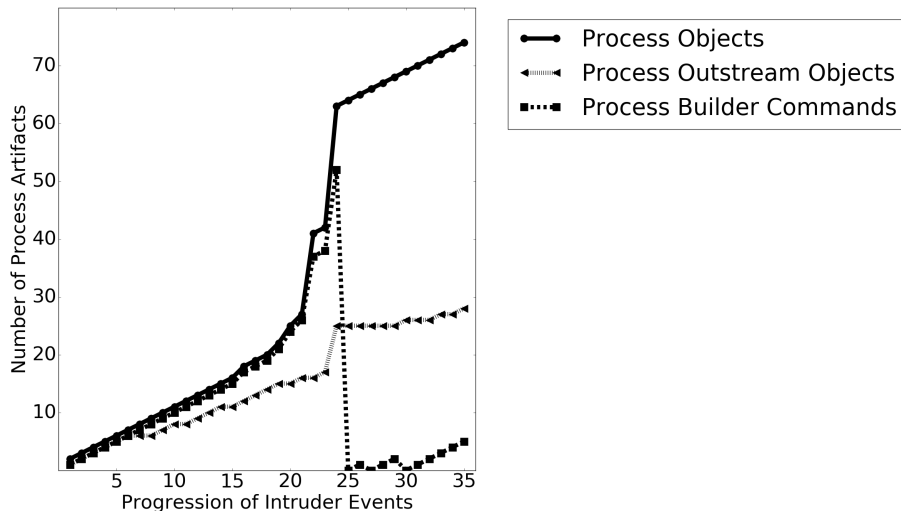
# Reconstructing Events

**Table:** This table shows a sampling of the processes started by the Java program and the `stdout` buffer at `t=21`.

| Address    | PID  | Buffered Data                    |
|------------|------|----------------------------------|
| 0x91dff7e0 | 1242 | #\n# This file MUST be edited w  |
| 0x91e1c7e8 | 1245 | Linux java-workx32-00 3.19.0-1   |
| 0x91e3b0e0 | 1248 | java adm cdrom sudo dip plugde   |
| 0x91e4a6e8 | 1250 | root:x:0:0:root:/root:/bin/bas   |
| 0x91eb1390 | 1252 | root:!:16678:0:99999:7::\ndaem   |
| 0x91f66708 | 1275 | \nStarting Nmap 6.47 ( http://n  |
| 0x91ff7ed0 | 1301 | history   grep pg\n history   gr |
| 0x92014f30 | 1307 | ifconfig\nsudo add-apt-reposito  |
| 0x920626d8 | 1322 | adding: home/java/.ssh/ (sto     |

# Evaluating Event Reconstruction: Process Objects

## RICE



# Conclusions

- Improves on existing memory analysis techniques
- Developed a methodology for recovering Java Objects
- Minimal solution achieves cross platform recovery
- Able to perform timelining and event reconstruction

# Conclusions

- Quantified these effects using TLS keys
- Developed a solution reducing latent secrets
- Developed a framework for recovering latent object
- Showed that latent objects can be used for forensics



- [1] N. Byte, “What the Heck was Stuxnet!?,” January 2016.
- [2] A. Drozhzhin, “Adwind malware-as-a-service hits more than 400,000 users globally,” February 2016.
- [3] V. Kamluk and A. Gostev, “ADWIND: A CROSS-PLATFORM RAT,” February 2016.
- [4] J. Scott-Railton, M. Marquis-Boire, C. Guarnieri, and M. Marschalek, “Packrat: Seven years of a south american threat actor,” December 2015.
- [5] V. Kamluk, C. Raiu, and I. Soumenkov, “THE â€œICEFOGâ€ APT: A TALE OF CLOAK AND THREE DAGGERS,” September 2013.
- [6] V. Kamluk, C. Raiu, and I. Soumenkov, “The icefog apt hits us targets with java backdoor,” January 2014.
- [7] E. Galperin, C. Quintin, M. Marquis-Boire, and C. Guarnieri, “I Got a Letter From the Government the



Other Day...:Unveiling a Campaign of Intimidation, Kidnapping, and Malware in Kazakhstan,” 2016.

- [8] J. Drake, “Exploiting memory corruption vulnerabilities in the java runtime,” 2011.
- [9] J. Viega, “Protecting sensitive data in memory,” 2001.
- [10] J. Chow, B. Pfaff, T. Garfinkel, K. Christopher, and M. Rosenblum, “Understanding data lifetime via whole system simulation,” in *Proceedings of the 13th Conference on USENIX Security Symposium - Volume 13*, SSYM’04, 2004.
- [11] J. Chow, B. Pfaff, T. Garfinkel, and M. Rosenblum, “Shredding your garbage: Reducing data lifetime through secure deallocation,” in *Proceedings of the 14th Conference on USENIX Security Symposium - Volume 14*, SSYM’05, 2005.

- [12] Y. Tang, P. Ames, S. Bhamidipati, A. Bijlani, R. Geambasu, and N. Sarda, “Cleanos: limiting mobile data exposure with idle eviction,” in *Presented as part of the 10th USENIX Symposium on Operating Systems Design and Implementation (OSDI 12)*, 2012.
- [13] M. Anikeev, F. C. Freiling, J. Götzfried, and T. Müller, “Secure garbage collection: Preventing malicious data harvesting from deallocated java objects inside the Dalvik VM,” *Journal of Information Security and Applications*, vol. 22, 2015.
- [14] Y. Li, “Where in your ram is “*python san\_diego.py*”?”, 2015.
- [15] A. Walters, “The Volatility Framework: Volatile memory artifact extraction utility framework,” 2007.
- [16] M. Cohen, “Rekall memory forensics framework,” in *DFIR Prague 2014*, SANS DFIR, 2014.

- [17] Z. Lin, X. Zhang, and D. Xu, “Automatic reverse engineering of data structures from binary execution,” 2010.
- [18] Z. Lin, J. Rhee, X. Zhang, D. Xu, and X. Jiang, “Siggraph: Brute force scanning of kernel data structure instances using graph-based signatures.,” in *NDSS*, 2011.
- [19] B. Dolan-Gavitt, A. Srivastava, P. Traynor, and J. Giffin, “Robust signatures for kernel data structures,” in *Proceedings of the 16th ACM conference on Computer and communications security*, ACM, 2009.
- [20] H. Macht, “Live memory forensics on Android with volatility,” *Friedrich-Alexander University Erlangen-Nuremberg*, 2013.
- [21] A. P. Heriyanto, “Procedures and tools for acquisition and analysis of volatile memory on Android smartphones,” 2013.

- [22] J. Park, H. Chung, and S. Lee, “Forensic analysis techniques for fragmented flash memory pages in smartphones,” *Digital Investigation*, vol. 9, no. 2, 2012.
- [23] A. Case, “Memory analysis of the Dalvik (Android) virtual machine,” Dec. 2011.
- [24] J. Okolica and G. Peterson, “Extracting the windows clipboard from memory,” in *Proceedings of the 2011 DFRWS Conference*, 2011.
- [25] A. Schuster, “Searching for processes and threads in microsoft windows memory dumps,” in *Proceedings of the 2006 DFRWS Conference*, 2006.
- [26] G. G. Richard III and V. Roussev, “Scalpel: A frugal, high performance file carver,” in *DFRWS*, 2005.

- [27] R. Beverly, S. Garfinkel, and G. Cardwell, “Forensic carving of network packets and associated data structures,” *digital investigation*, vol. 8, 2011.
- [28] S. Hand, Z. Lin, G. Gu, and B. Thuraisingham, “Bin-carver: Automatic recovery of binary executable files,” *Digital Investigation*, vol. 9, 2012.
- [29] A. Pridgen, S. Garfinkel, and D. S. Wallach, “Present but unreachable: Reducing persistent latent secrets in HotSpot JVM,” in *System Science (HICSS), 2017 50th Hawaii International Conference on*, IEEE, 2017.